

## Lab 12: Multithreading

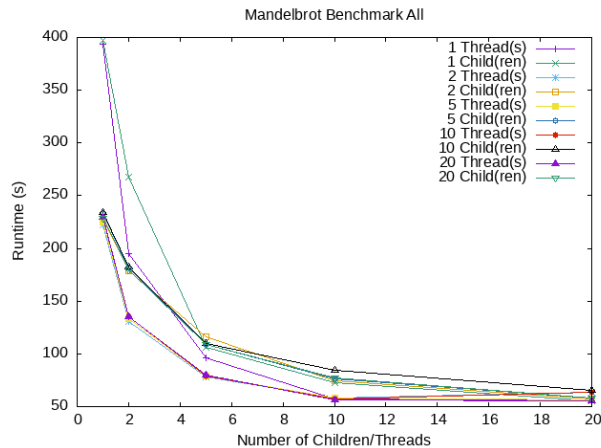
### Overview of Threads

I created a struct to store the arguments for each thread, and created a method for each thread to execute to build the image.

I used the 'pthread' library in C to create and manage threads. The main thread waited for all child threads to complete their execution using the 'pthread\_join' function.

## Results

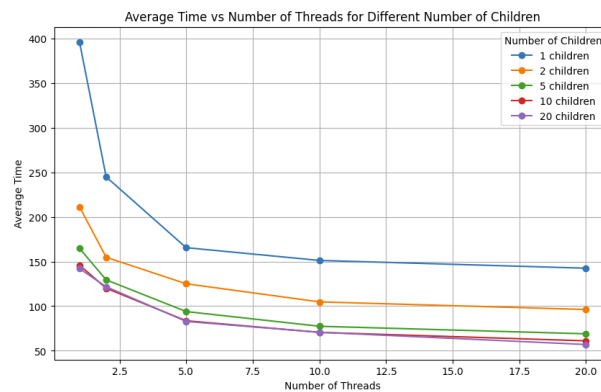

num_children	num_threads	time
1	1	394.052345
2	1	194.664908
5	1	96.58268
10	1	57.393608
20	1	55.140141
1	2	222.368145
2	2	130.764116
5	2	77.738222
10	2	57.973349
20	2	63.703118
1	5	224.807174
2	5	133.744711
5	5	79.024139
10	5	57.813655
20	5	56.56309
1	10	230.097656
2	10	135.418242
5	10	79.059961
10	10	57.041502
20	10	63.883061
1	20	229.269893
2	20	135.426675
5	20	79.54736
10	20	56.744055
20	20	55.613426



Multithreading seemed to impact runtime more. This is likely because the overhead of creating and managing threads is less than the overhead of creating and managing processes.

- Was there a “sweet spot” where optimal (minimal) runtime was achieved?

Yes, the optimal runtime was achieved with 20 threads and 10 children.



## Discussion

- Which technique seemed to impact runtime more – multithreading or multiprocessing. Why do you think that is?